

Get to know CouchDB

My struggle:

My struggle:

- Oh, look another NoSQL DB!

My struggle:

- Oh, look another NoSQL DB!
- I know HTTP!

My struggle:

- Oh, look another NoSQL DB!
- I know HTTP!
- -> CouchDB should easy

My struggle:

- Oh, look another NoSQL DB!
- I know HTTP!
- -> CouchDB should easy
- Weeks later still **no clue how it works** :(

My struggle:

- Oh, look another NoSQL DB!
- I know HTTP!
- -> CouchDB should easy
- Weeks later still **no clue how it works** :(
- Visited CouchDB Meetup / CouchDB Day

My struggle:

- Oh, look another NoSQL DB!
- I know HTTP!
- -> CouchDB should easy
- Weeks later still **no clue how it works** :(
- Visited CouchDB Meetup / CouchDB Day
- After *months* i grasped the concept

My struggle:

- Oh, look another NoSQL DB!
- I know HTTP!
- -> CouchDB should easy
- Weeks later still **no clue how it works** :(
- Visited CouchDB Meetup / CouchDB Day
- After *months* i grasped the concept
- Able to use it now :)

My struggle:

- Oh, look another NoSQL DB!
- I know HTTP!
- -> CouchDB should easy
- Weeks later still **no clue how it works** :(
- Visited CouchDB Meetup / CouchDB Day
- After *months* i grasped the concept
- Able to use it now :)

My struggle:

- Oh, look another NoSQL DB!
- I know HTTP!
- -> CouchDB should easy
- Weeks later still **no clue how it works** :(
- Visited CouchDB Meetup / CouchDB Day
- After *months* i grasped the concept
- Abled to use it now :)

Let me tell you how it works in ~40 Minutes!

- What is CouchDB?
- Simple usage
 - Databases, Documents, CRUD
- How CouchDB works
 - B-tree, ranges, pagination
- Advanced usage
 - Views, Grouping
- Extended topics
 - More fancy stuff if time left...

Julius Beckmann

Silpion IT-Solutions

- PHP / Symfony2
- Erlang / Elixir
- NodeJS
- Continuous Everything!

What is CouchDB?

What is CouchDB?

- **Cluster Of Unreliable Commodity Hardware DB**
- Started 2005 by [@damienkatz](#) at Lotus Notes.
- Since 2008 Apache Project.
- Open Source: Apache Software license 2.0.

- RESTful HTTP as Interface.
- Schemaless JSON documents.
- Written in Erlang.
- Version 1.6 is current.
- Version 2.0 with *real clustering* on the way!



www.meetup.com/CouchDB-Meetup-Hamburg

Organized by

- Robert Kowalski @robinson_k
- Andreas Wenk @awenkh
- Klaus Trainer @KlausTrainer

Installation

Webinterface: http://localhost:5984/_utils/.

Debian/Ubuntu

```
apt-get install couchdb
```

<http://wiki.ubuntuusers.de/couchdb>

Docker image

Using this Image: [klaemo/docker-couchdb](#)

- `docker pull klaemo/couchdb:latest`
- `docker run -d -p 5984:5984 --name couchdb klaemo/couchdb`

Stop and start container:

- `docker stop couchdb`
- `docker start couchdb`

Simple usage

Say hello to CouchDB

```
$ curl -X GET http://localhost:5984/
```

```
{  
  "couchdb": "Welcome",  
  "uuid": "460a0cc31a109fe7cdcdd17f2109e515",  
  "version": "1.6.1",  
  "vendor": {  
    "version": "1.6.1",  
    "name": "The Apache Software Foundation"  
  }  
}
```

List Databases

```
$ curl -X GET http://localhost:5984/_all_dbs
```

```
[  
  "_replicator",  
  "_users"  
]
```

- ① What is a Database
- ② What is a Document
- ③ CRUD for Documents
- ④ CRUD for Attachments

1. What is a Database

- Single container = No tables.
- Permissions granted on database level by CouchDB.

1. What is a Database

- Single container = No tables.
- Permissions granted on database level by CouchDB.

Create Database

```
$ curl -X PUT http://localhost:5984/my_database
```

```
{"ok":true}
```


1. What is a Database

- Single container = No tables.
- Permissions granted on database level by CouchDB.

Create Database

```
$ curl -X PUT http://localhost:5984/my_database
```

```
{"ok":true}
```

Remove Database

```
$ curl -X DELETE http://localhost:5984/my_database
```

```
{"ok":true}
```

Database Metadata

```
$ curl -X GET http://localhost:5984/my_database
```

```
{
  "db_name": "my_database",
  "doc_count": 0,
  "doc_del_count": 0,
  "update_seq": 0,
  "purge_seq": 0,
  "compact_running": false,
  "disk_size": 79,
  "data_size": 0,
  "instance_start_time": "1441012193869017",
  "disk_format_version": 6,
  "committed_update_seq": 0
}
```

2. What are Documents

- JSON-Format.
- Unicode charset.
- Schemaless structure.
- Underscore prefix is reserved for first level keys:
 - `_id`
 - `_rev`
 - `_attachments`
 - `_deleted`

Creating a Document

```
// test.json
{
  "hello": "world",
  "age": 42,
  "tags": ["couchdb", "is", "cool"]
}
```

Creating a Document

```
// test.json
{
  "hello": "world",
  "age": 42,
  "tags": ["couchdb", "is", "cool"]
}
```

```
$ curl -X POST http://localhost:5984/my_database
  -H "Content-Type: application/json"
  --data-binary @test.json
```

```
{
  "ok": true,
  "id": "22d26e752e31bd1572f1d20d94000c47",
  "rev": "1-ff73263039516726e7917f24111c80ec"
}
```

- By default UUIDs.
- Must be unicode **string**.
- Will be part of URL.
- Can be chosen for each document.

CouchDB IDs

- By default UUIDs.
- Must be unicode **string**.
- Will be part of URL.
- Can be chosen for each document.

```
$ curl -X PUT http://localhost:5984/my_database/hello-world  
-H "Content-Type: application/json"  
--data-binary @test.json
```

```
{  
  "ok":true,  
  "id":"hello-world",  
  "rev":"1-ff73263039516726e7917f24111c80ec"  
}
```

- *Multi-Version Concurrency Control.*
- 'Strictly monotonic ascending'.
- Revision of document.

- *Multi-Version Concurrency Control.*
- 'Strictly monotonic ascending'.
- Revision of document.
- Optimistic locking => lock-free access.
- Needed for replication.
- Initial $i = 1$.

```
rev = i + "-" + md5(body,attachments,deleted)  
  
"1-ff73263039516726e7917f24111c80ec"
```

Fetching a Document

*Skipping **http://localhost:5984** in url from now on...*

```
$ curl -X GET /my_database/hello-world
```

```
{  
  "_id": "hello-world",  
  "_rev": "1-ff73263039516726e7917f24111c80ec",  
  "hello": "world",  
  "age": 42,  
  "tags": ["couchdb", "is", "cool"]  
}
```

Updating a Document

```
// test_update.json
{
  "hello": "world!",
  "age": 43,
  "tags": ["couchdb", "is", "cooler"]
}
```

Updating a Document

```
// test_update.json
{
  "hello": "world!",
  "age": 43,
  "tags": ["couchdb", "is", "cooler"]
}
```

```
$ curl -X PUT /my_database/hello-world?rev=1-ff732630...
  -H "Content-Type: application/json"
  --data-binary @test_update.json
```

```
{
  "ok": true,
  "id": "hello-world",
  "rev": "2-b9f6bbfa08e0f6a8cf3d0dcd4c153fc7"
}
```

Deleting a Document

```
$ curl -X DELETE /my_database/hello-world?rev=2-b9f6bbfa...
```

```
{  
  "ok":true,  
  "id":"hello-world",  
  "rev":"3-e4f06ea31419b24495c94bcd952777b6"  
}
```

Deleting a Document

```
$ curl -X DELETE /my_database/hello-world?rev=2-b9f6bbfa...
```

```
{
  "ok":true,
  "id":"hello-world",
  "rev":"3-e4f06ea31419b24495c94bcd952777b6"
}
```

```
$ curl -X GET http://localhost:5984/my_database/hello-world
{"error":"not_found","reason":"deleted"}
```

- Document attribute `_deleted`: `true`
- *OLD* revisions still there till **compact**:

```
$ curl -X GET /my_database/hello-world?rev=2-b9f6bbfa...
```

- Files attached to documents
- Binary data of any kind: Images, PDF, Mails ...

API

- **PUT** /products/apple-macbook-air/preview.jpg?rev=XXX
- **GET** /products/apple-macbook-air/preview.jpg
- **DELETE** /products/apple-macbook-air/preview.jpg?rev=XXX


Insights we learned:

- ➊ Documents have `_id`'s and `_rev`'s.
- ➋ Update/Delete needs latest `_rev`.
- ➌ Optimistic locking via `_rev`.
- ➍ Update **replaces** document.
- ➎ Binary attachments.

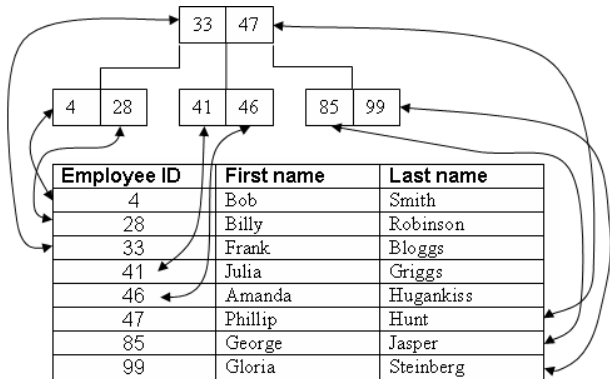
How CouchDB works

Databases use Trees to be efficient

Databases use Trees to be efficient

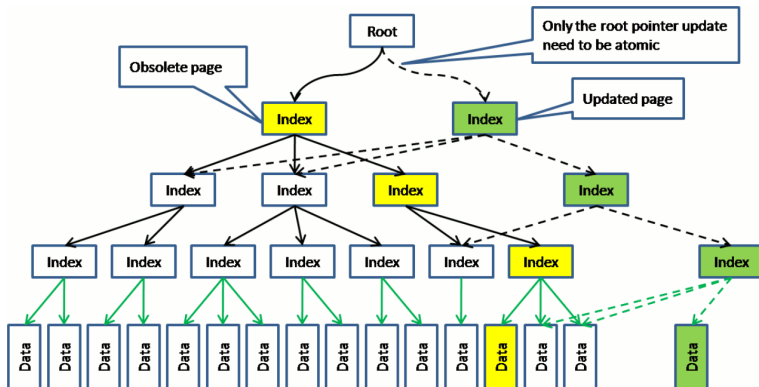


Databases use Trees to be efficient



CouchDB is using a B+ Tree

Understanding the B-Tree



- Using B-Tree to access documents.
- Efficient low tree structure.

IDs are unicode *strings*!

IDs are unicode *strings*!

There is a order defined by their binary value:

- A = U0041 to Z = U005A
- First: U0000 - Last: UFFFF

IDs are unicode *strings*!

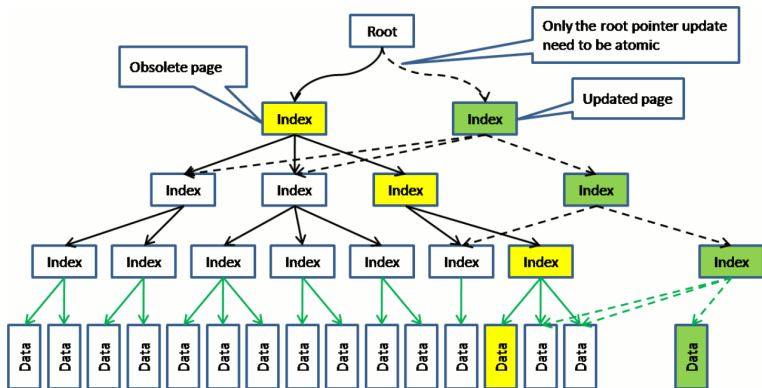
There is a order defined by their binary value:

- A = U0041 to Z = U005A
- First: U0000 - Last: UFFFF

IDs are held in *index* of Tree, not in *data*.

=> **Range iteration** possible.

Understanding the B-Tree



- From U0000 on the *left*
- To UFFFF on the *right*

Documents

```
{ "_id": "lenovo-thinkpad-t440", ... }  
{ "_id": "lenovo-edge-e130", ... }  
{ "_id": "apple-macbook-2013", ... }  
{ "_id": "apple-macbook-pro-2015", ... }
```

Usefull IDs

Documents

```
{ "_id": "lenovo-thinkpad-t440", ... }  
{ "_id": "lenovo-edge-e130", ... }  
{ "_id": "apple-macbook-2013", ... }  
{ "_id": "apple-macbook-pro-2015", ... }
```

“Query” all apple products

```
$ curl '/products/_all_docs?startkey="apple-"\n      &endkey="apple-\\uffff"'
```

- `_all_docs` is the resource of the *whole B-Tree* of the database.

CouchDB is a restfull HTTP API for B+ Trees!

Typical CouchDB IDs

- user-alice, user-bob
- product-apple-macbook-2013, product-lenovo-thinkpad-t440
- order-user-alice-20150102, order-user-bob-20141212

Typical CouchDB IDs

- user-alice, user-bob
- product-apple-macbook-2013, product-lenovo-thinkpad-t440
- order-user-alice-20150102, order-user-bob-20141212

All users

```
$ curl '/my_database/_all_docs?startkey="user-"\
      &endkey="user-\\uffff"'
```

Typical CouchDB IDs

- user-alice, user-bob
- product-apple-macbook-2013, product-lenovo-thinkpad-t440
- order-user-alice-20150102, order-user-bob-20141212

All users

```
$ curl '/my_database/_all_docs?startkey="user-"\&endkey="user-\\uffff"'
```

All products by lenovo

```
$ curl '/my_database/_all_docs?startkey="product-lenovo-"\&endkey="product-lenovo-\\uffff"'
```

Typical CouchDB IDs

- user-alice, user-bob
- product-apple-macbook-2013, product-lenovo-thinkpad-t440
- order-user-alice-20150102, order-user-bob-20141212

All users

```
$ curl '/my_database/_all_docs?startkey="user-"\
      &endkey="user-\\uffff"'
```

All products by lenovo

```
$ curl '/my_database/_all_docs?startkey="product-lenovo-"\
      &endkey="product-lenovo-\\uffff"'
```

All orders by bob in 2014

```
$ curl '/my_database/_all_docs?startkey="order-user-bob-2014"\
      &endkey="order-user-bob-2014\\uffff"'
```


Pagination using Keys

Simple pagination

```
$ curl '/my_database/_all_docs?limit=10&skip=0'  
$ curl '/my_database/_all_docs?limit=10&skip=10'
```

Pagination using Keys

Simple pagination

```
$ curl '/my_database/_all_docs?limit=10&skip=0'  
$ curl '/my_database/_all_docs?limit=10&skip=10'
```

Paginate user documents

First page:

```
$ curl '/my_database/_all_docs?startkey="user-"\n      &endkey="user-\\uffff"&limit=10'
```

Second page via `skip`:

```
$ curl '/my_database/_all_docs?startkey="user-"\n      &endkey="user-\\uffff"&limit=10&skip=10'
```

Second page via `startkey_docid`:

```
$ curl '/my_database/_all_docs?startkey_docid=user-alice\n      &endkey="user-\\uffff"&limit=10'
```

Fetching multiple keys at once

```
// keys.json
{
  "keys" : [
    "user-alice",
    "order-user-alice-20150102"
  ]
}
```

```
$ curl -X POST /my_database/_all_docs
-H "Content-Type: application/json"
--data-binary @keys.json
```

Insights we learned:

- ➊ Restfull HTTP API to B+ Tree.
- ➋ Meaningfull IDs instead of UUIDs.
- ➌ IDs can determine type of a document.
- ➍ Iteration over IDs is usefull.
- ➎ Combination of keys is very usefull.
- ➏ Dates in IDs can be handy.

But this is not enough to work with ...

We need:
queries!
grouping!
database logic!

Advanced usage

Querys **always** need a Index
and a Index is a Tree.

Querys **always** need a Index
and a Index is a Tree.

We have build that index ourself.

Views

CouchDB offers **Map**-Reduce to create indexes.

Map function

```
function(doc) {  
    //    key        value  
    emit(doc.age, doc.hello);  
}
```

Result

- Key: 42, Value: "world"
- Key: 43, Value: "world!"

CouchDB offers Map-**Reduce** to create indexes.

Reduce function

```
function (key, values, rereduce) {  
    return values.join(',');  
}
```

Result

- "world,world!"

Map Functions:

- Each map function called *once* for any changed doc
- Define key and value of index
- Convert data to needed format
- Filter documents
- Create multiple entries from one doc

Reduce Functions:

- Will be called on collection of mapped entries
- Included functions: `_sum`, `_count`, `_stats`
- *Summarizes* mapped data

Example: Query for Tags

Map function:

```
function(doc) {  
    doc.tags.forEach(function(tag) {  
        emit(tag, doc._id);  
    });  
}
```

Result:

```
"cool": "22d26e752e31bd1572f1d20d94000c47"  
"cooler": "hello-world",  
"couchdb": "22d26e752e31bd1572f1d20d94000c47"  
"couchdb": "hello-world"  
"is": "22d26e752e31bd1572f1d20d94000c47"  
"is": "hello-world"
```

Reduce function:

```
_count
```

Result with reduce:

```
"cool": 1
```

```
"cooler": 1
```

```
"couchdb": 2
```

```
"is": 2
```

Example: Index date ranges

Example for flexibility of CouchDB views

Map function:

```
function(doc) {  
  
    var dates = date_range(doc.start, doc.end);  
    // ['2015-04-01', '2015-04-02', '2015-04-03', ...]  
  
    dates.forEach(function(date) {  
        emit(date, doc._id);  
    });  
}
```

- Index will contain data, that not existed in doc!

Design Documents

- JSON document
- Collection of javascript code
 - for views
 - for callbacks
- Container for “application code”
- Multiple Design Documents common
 - for users
 - for products
 - ...
- *Short: **DDoc***

Design Documents

```
{
  "_id": "_design/tags",
  "_rev": "1-753abddc33070499276277bd90decabd",
  "language": "javascript",
  "views": {
    "by-tag": {
      "map": "function(doc) {
                doc.tags.forEach(function(tag) {
                  emit(tag, doc._id);
                });
              }",
      "reduce": "_count"
    }
  }
}
```

Access the *by-tags* view from ddoc *tags*:

```
$ curl -X GET '/my_database/_design/tags/_view/by-tag'
```

Access the *by-tags* view from ddoc tags:

```
$ curl -X GET '/my_database/_design/tags/_view/by-tag'
```

```
{  
  "rows": [  
    {  
      "key": null,  
      "value": 6  
    }  
  ]  
}
```

Whats that? - Ahh, its reduced ...

Query View Reduce Flag

```
$ curl -X GET '/my_database/_design/tags/_view/by-tag\
?reduce=false'
```

```
{
  "offset": 0,
  "total_rows": 6,
  "rows": [
    {
      "id": "22d26e752e31bd1572f1d20d94000c47",
      "key": "cool",
      "value": "22d26e752e31bd1572f1d20d94000c47"
    },
    {
      "id": "hello-world",
      "key": "cooler",
      "value": "hello-world"
    },
    ...
  ]
}
```

Query View Group Level

```
$ curl -X GET '/my_database/_design/tags/_view/by-tag\
?group_level=0'
```

```
{
  "rows": [
    {
      "key": null,
      "value": 6
    }
  ]
}
```

The same as the **first request**.

Default is ?group_level=0
and reduce=true if reduce function exists.

Query View Group Level

```
$ curl -X GET '/my_database/_design/tags/_view/by-tag\  
?group_level=1'
```

```
{  
  "rows": [  
    {  
      "key": "cool",  
      "value": 1  
    },  
    {  
      "key": "couchdb",  
      "value": 2  
    },  
    ...  
  ]  
}
```


Query View by Key

```
$ curl -X GET '/my_database/_design/tags/_view/by-tag\
?key="couchdb"&reduce=false'
```

```
{
  "offset": 2,
  "total_rows": 6,
  "rows": [
    {
      "id": "22d26e752e31bd1572f1d20d94000c47",
      "key": "couchdb",
      "value": "22d26e752e31bd1572f1d20d94000c47"
    },
    {
      "id": "hello-world",
      "key": "couchdb",
      "value": "hello-world"
    }
  ]
}
```

Insights we learned:

- 1 Views instead of Queries.
- 2 Views programmed with **map-reduce**.
- 3 Each view is a own accessible Index.
- 4 Code for views inside a Design-Document.
- 5 No arbitrary WHERE `this=$that AND`

Extended Topics

Different documents in same database

- Using prefix for id
- Having type field

Simply using both:

```
{  
  "_id": "user-alice",    // for _all_docs  
  "_rev": "1-753abddc33070499276277bd90decabd",  
  "type": "user",        // for views  
  "name": "Alice"  
}
```

- Can contain multiple views and other functions.
- callbacks:
 - views
 - filters
 - lists
 - rewrites
 - shows
 - updates
 - validate_doc_update
- In Javascript, Ruby, Python, Perl, Lisp, Erlang, ...
- Changing DDoc = rebuilding views.
- Building views costs time!

Changes feed

- Like twitter for your database, with less drama.
- Easy notification system.

```
$ curl /my_database/_changes?since=5
```

```
{
  "last_seq":6,
  "results":[
    {
      "changes":[{"rev":"3-e4f06ea31419b24495c94bcd952777b6"}],
      "deleted":true,
      "id":"hello-world",
      "seq":6
    }
  ]
}
```

Replication

CouchDB can **replicate DB** from other CouchDB via HTTP API.

- 1 Master - n Slaves
- n Master
- ...
- Slaves of Slaves possible
- *“Automatic conflict handling”* ...

CAP

For CouchDB in *Multi-Master cluster* Setup.

CAP

- Partition Tolerance: YES
- Availability: YES
- Consistency: Eventual

Pick Two

- CA: **MySQL**, Postgres, RDBMS
- CP: **MongoDB**, Google BigTable
- AP: **CouchDB**, Cassandra

- CouchDB fits best for Read >> Write.

- CouchDB fits best for Read >> Write.
- Use HTTP correctly (ETag, ...)

- CouchDB fits best for Read >> Write.
- Use HTTP correctly (ETag, ...)
- Create new docs instead of updating.

- CouchDB fits best for Read >> Write.
- Use HTTP correctly (ETag, ...)
- Create new docs instead of updating.
- Need to design to avoid conflicts.

- CouchDB fits best for Read >> Write.
- Use HTTP correctly (ETag, ...)
- Create new docs instead of updating.
- Need to design to avoid conflicts.
- Design for your access pattern.

- CouchDB fits best for Read >> Write.
- Use HTTP correctly (ETag, ...)
- Create new docs instead of updating.
- Need to design to avoid conflicts.
- Design for your access pattern.
- Avoid highly relational data.

- CouchDB fits best for Read >> Write.
- Use HTTP correctly (ETag, ...)
- Create new docs instead of updating.
- Need to design to avoid conflicts.
- Design for your access pattern.
- Avoid highly relational data.
- Small attachment files only.

- CouchDB fits best for Read >> Write.
- Use HTTP correctly (ETag, ...)
- Create new docs instead of updating.
- Need to design to avoid conflicts.
- Design for your access pattern.
- Avoid highly relational data.
- Small attachment files only.
- Doc revisions ARE NOT A Revision Control System!

- CouchDB fits best for Read >> Write.
- Use HTTP correctly (ETag, ...)
- Create new docs instead of updating.
- Need to design to avoid conflicts.
- Design for your access pattern.
- Avoid highly relational data.
- Small attachment files only.
- Doc revisions ARE NOT A Revision Control System!
- Reduce function `_sum` and `_count` are really cheap.

- Lots of disk space because of appending file format.
- Compaction of database/views.
 - Rewriting whole file.
 - Removing old `_rev`'s.
- Can be done automatically:
 - wiki.apache.org/couchdb/Compaction

Further Topics to explore

- Clustering
 - via HTTP Replication
 - Master-Slave
 - Master-Master
- Browser Local Storage
 - PouchDB
 - Syncing Browser of offline Apps
- Empowering HTTP
 - Proxy
 - Caching (Etag, ...)
- Fulltext search
 - River to Elasticsearch
 - Push to Lucene

Links

- couchdb.apache.org
- wiki.apache.org/couchdb
- guide.couchdb.org
- <http://docs.ehealthafrica.org/couchdb-best-practices/>

Other talks

- <https://speakerdeck.com/renatosnrg/couchdb>
- <https://speakerdeck.com/dzuelke/an-introduction-to-couchdb-ipc2011se-2011-06-01>

The End - Thanks!